

Tubi's Story with gRPC (in Scala)

Scala Meetup @ 2023/01/14

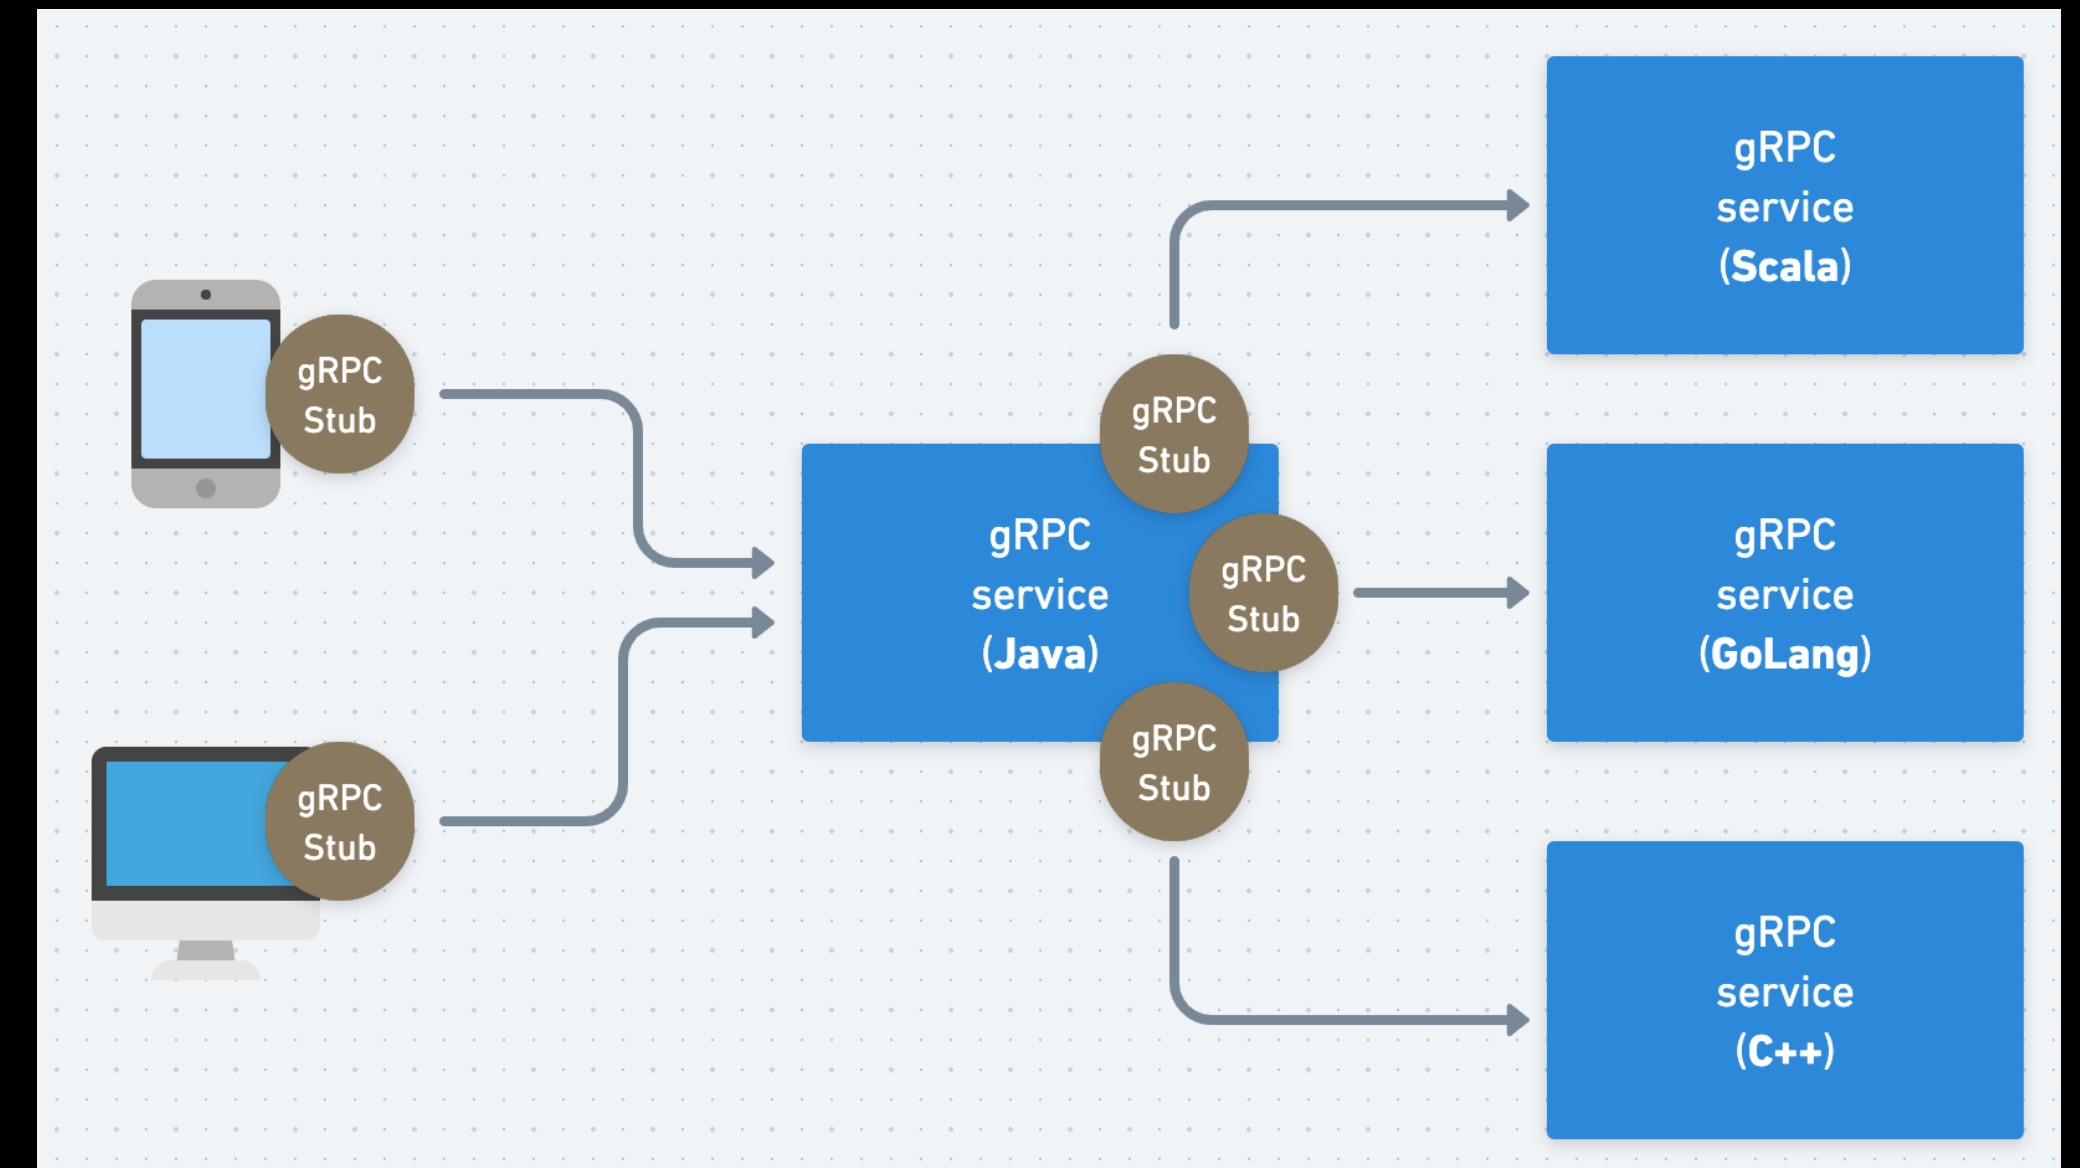
张正来 <zhenglaizhang@tubi.tv>

目录


- gRPC简介
- Tubi Scala team & gRPC
- Akka-gRPC 迁移
- Q&A


什么是gRPC


- 高性能、开源和通用的 RPC 框架
- Google (2015) => CNCF (2017)
- Cloud Native和Microservices系统通信
- 已被业内广泛应用
- g
 - **g**rpc **r**emote **p**rocedure **c**all
 - 'g' stands for something different





Used by


 Square

 NETFLIX


 CoreOS

 Cockroach LABS

 CISCO

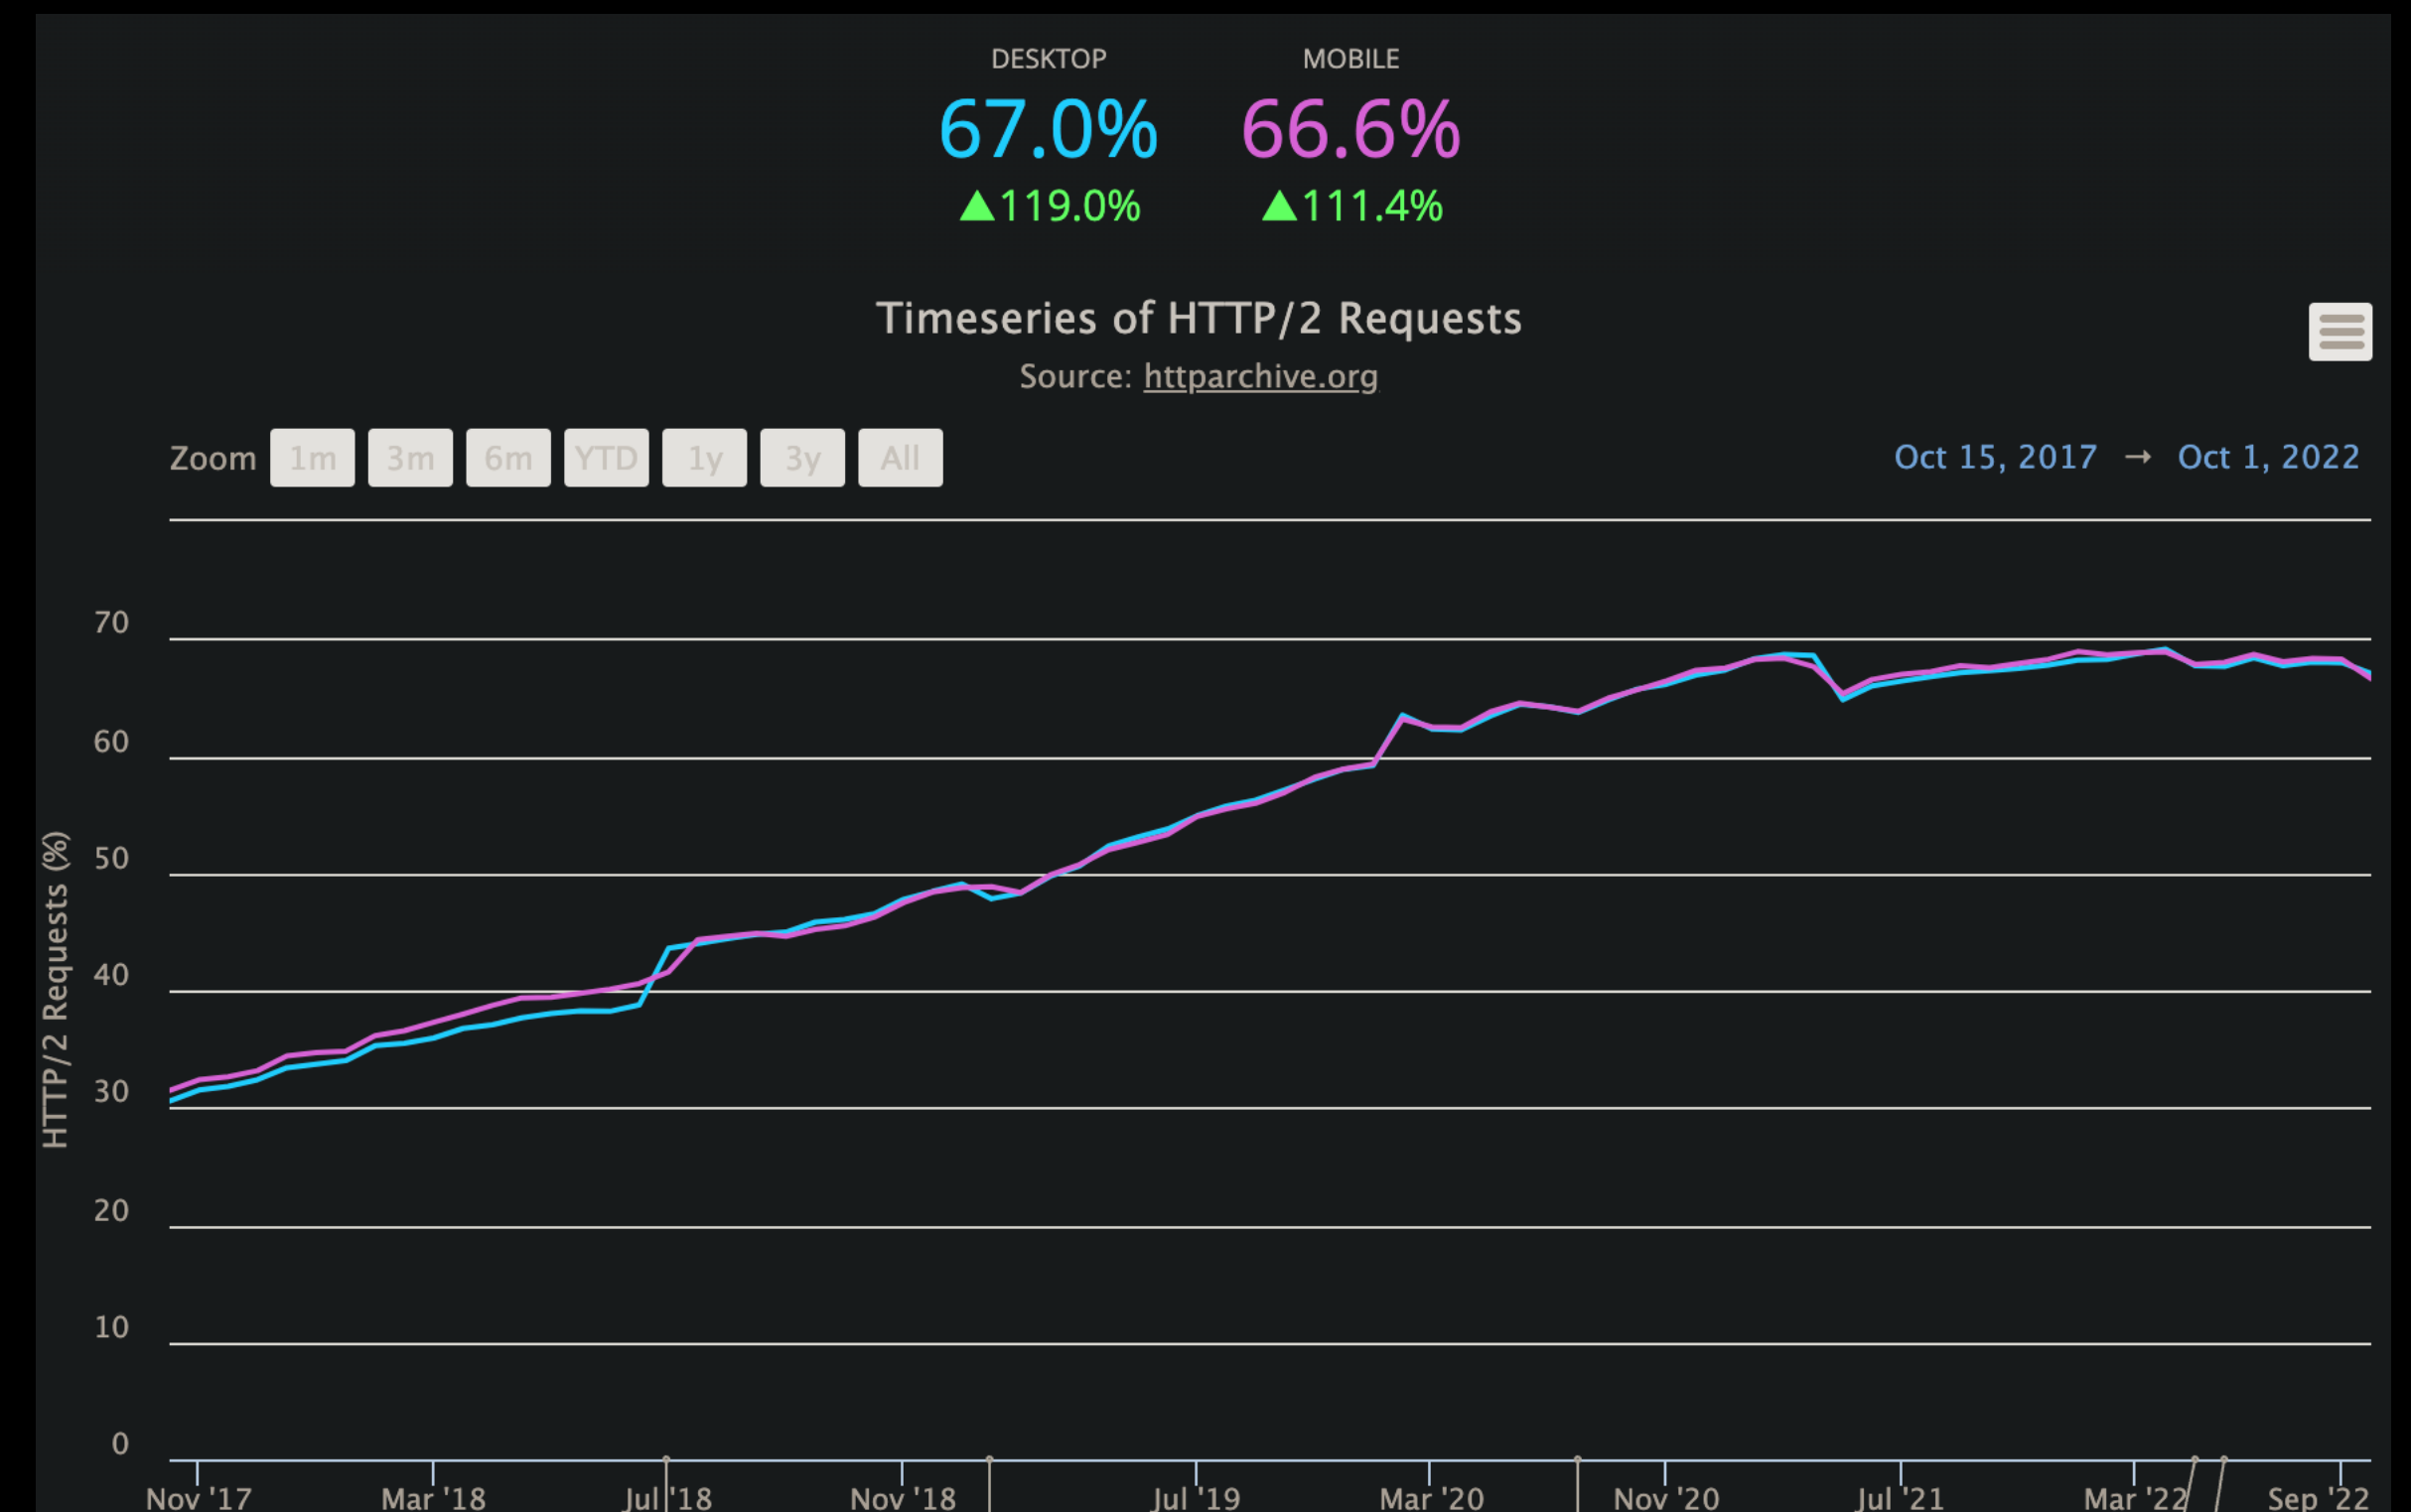
 JUNIPER NETWORKS

gRPC is a CNCF incubation project

 **CLOUD NATIVE**
COMPUTING FOUNDATION

gRPC特性 (一)

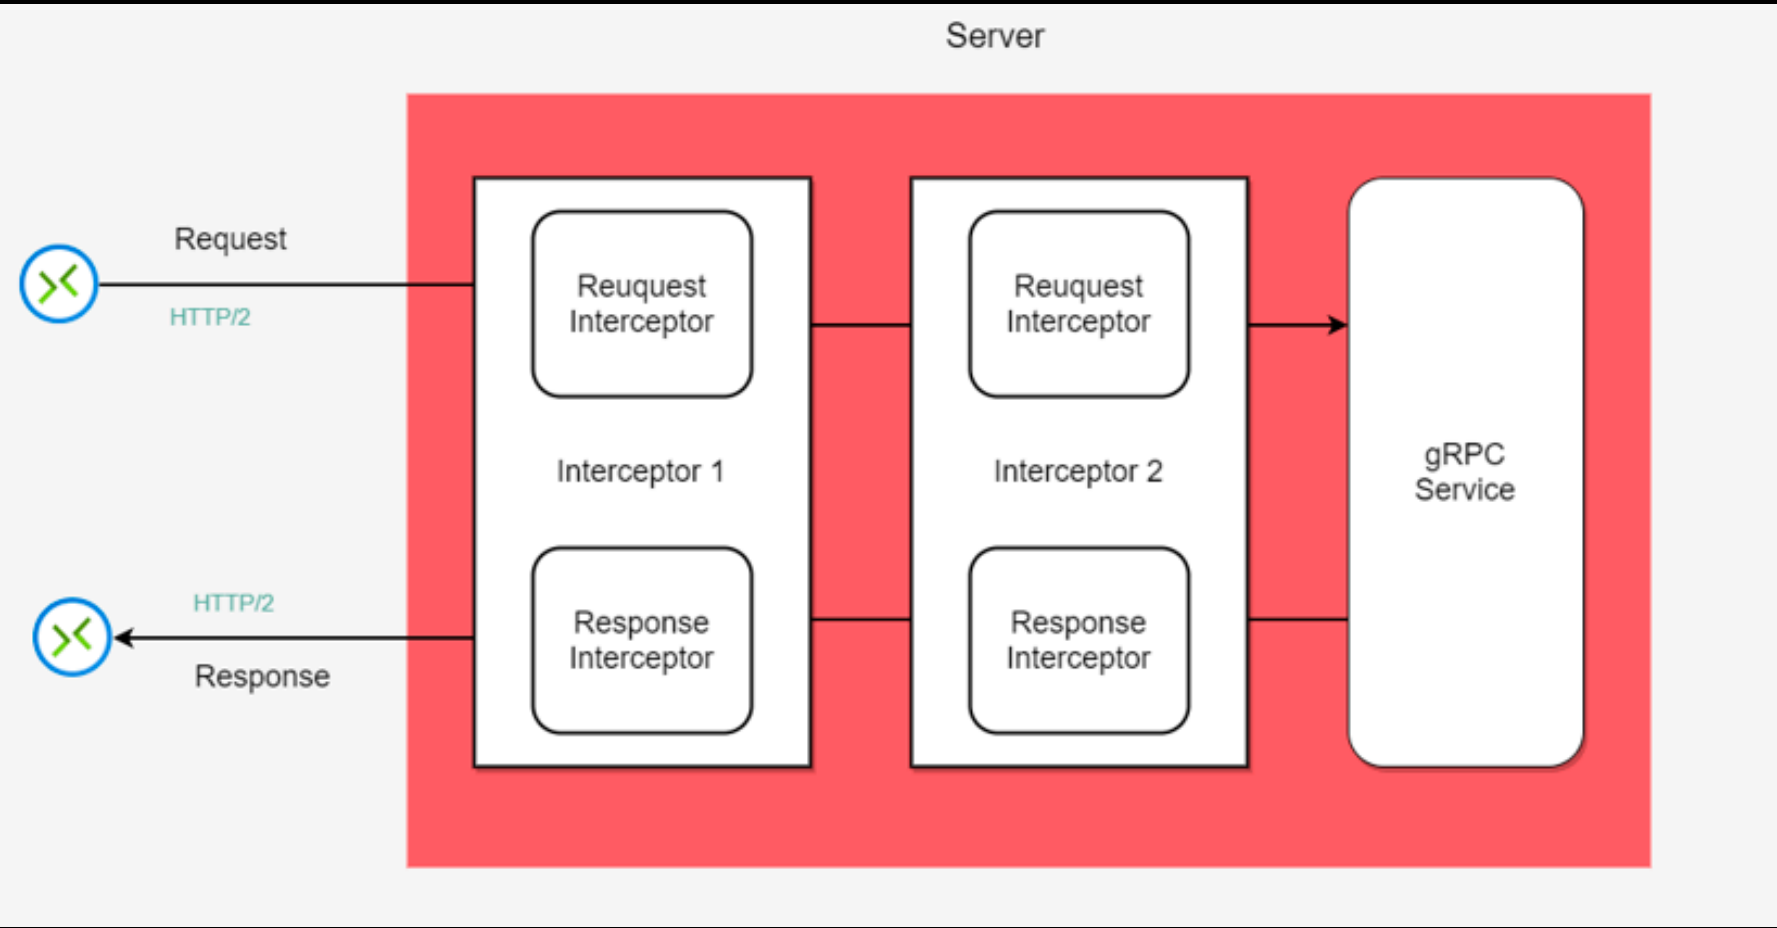
- 编码协议(IDL)默认基于Protocol Buffers
 - 接口优先
 - 高效
 - 可扩展
 - 兼容性好
- 传输协议基于 HTTP/2
 - 二进制帧
 - 流式传输和多路复用 (multiplexing)
 - 元数据 (metadata)
- 跨语言、跨平台
- 高性能



Ref: <https://httparchive.org/reports/state-of-the-web#h2>

gRPC特性（二）

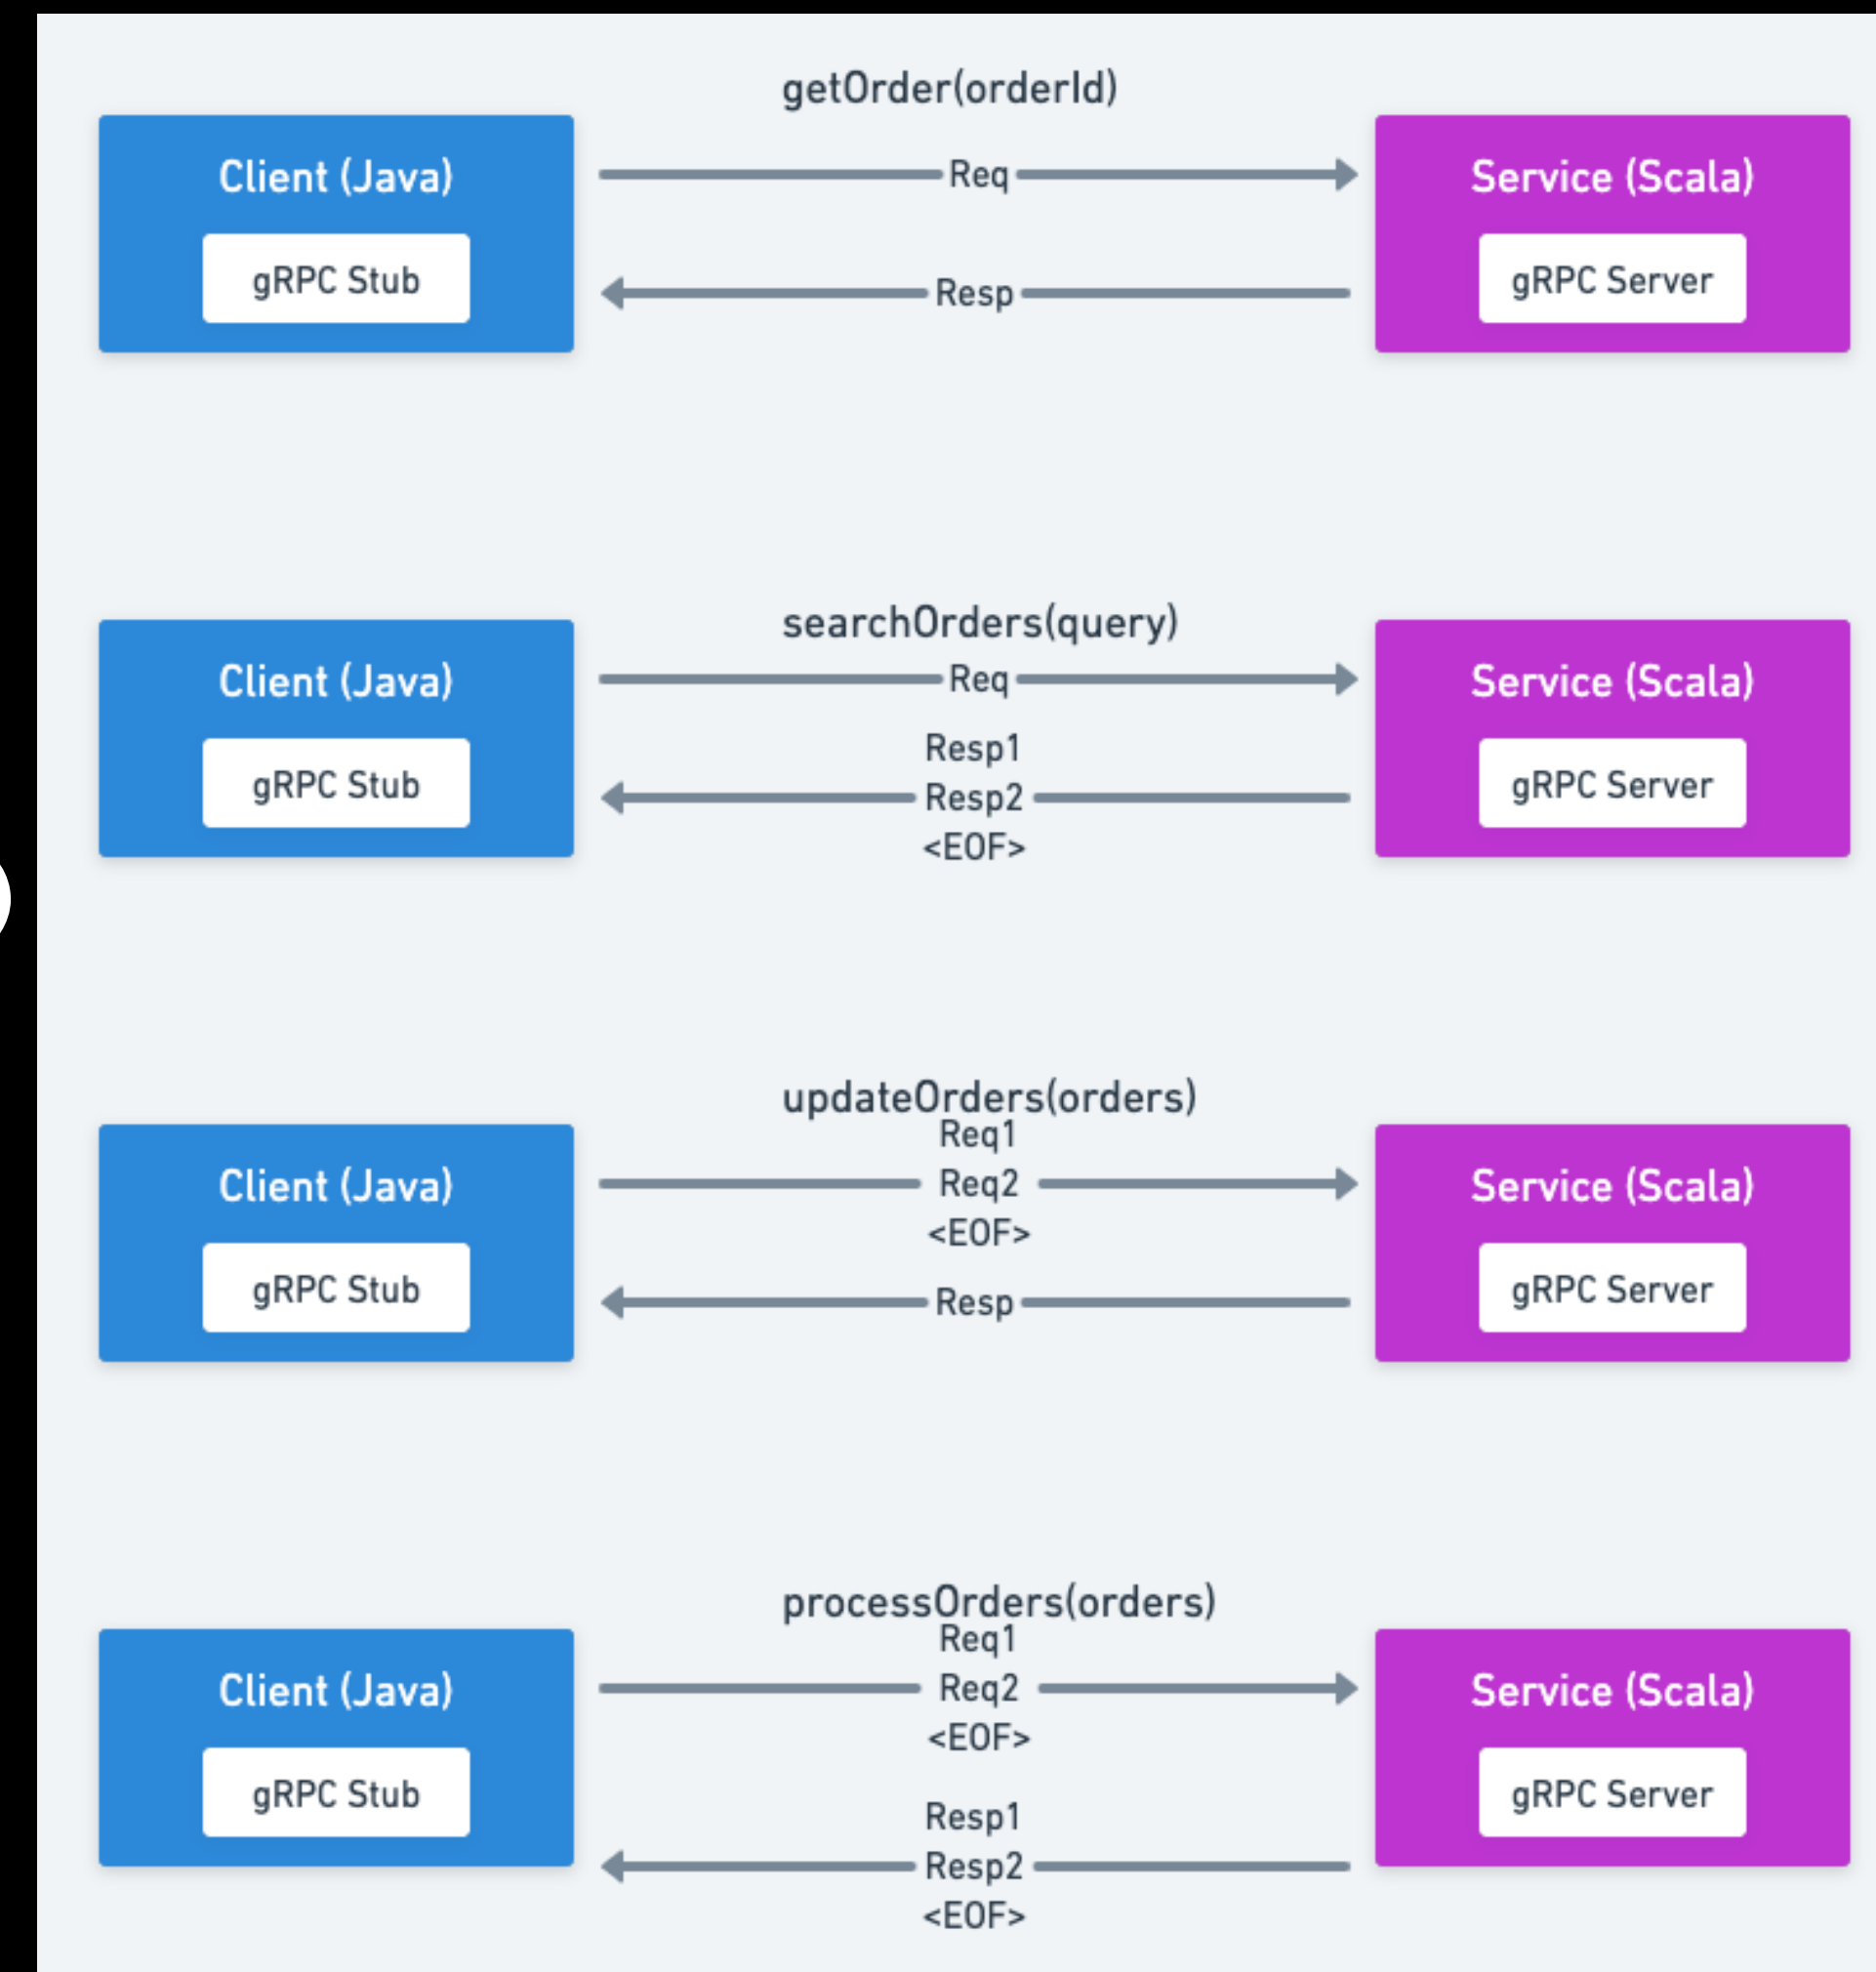
- 截止时间（deadline）
- 取消（cancellation）
- 错误处理
- 拦截器（interceptor）扩展机制
- Server reflection
- TLS加密



gRPC状态码	gRPC信息	含义
0	OK	成功
1	CANCELLED	操作已被取消
2	UNKNOWN	未知错误
3	INVALID_ARGUMENT	客户端参数非法
4	DEADLINE_EXCEEDED	操作超过了截止时间
5	NOT_FOUND	请求实体未找到
6	ALREADY_EXISTS	客户端试图创建的实体已存在
7	PERMISSION_DENIED	调用者没有权限执行特定操作
8	RESOURCE_EXHAUSTED	资源已耗尽
9	FAILED_PRECONDITION	操作被拒绝，系统没有处于执行操作所需状态
10	ABORTED	操作被中止
11	OUT_OF_RANGE	操作超出了合法的范围
12	UNIMPLEMENTED	该操作未实现
13	INTERNAL	内部错误
14	UNAVAILABLE	服务当前不可用
15	DATA_LOSS	数据丢失或损坏
16	UNAUTHENTICATED	客户端没有进行操作的合法认证凭证

gRPC通信模式(Demo)

- 单项RPC (unary/simple RPC)
- 服务端流式 RPC (server-side streaming RPC)
- 客户端流式 RPC (client-side streaming RPC)
- 双向流式 RPC (bidirectional streaming RPC)



框架对比

	SOAP	REST	Thrift	gRPC	GraphQL
发布时间	90年代末	2000	2007	2015	2015
格式	XML	JSON, XML, etc	JSON,Binary	默认ProtoBuf, JSON	JSON
传输协议	HTTP/TCP	HTTP	TCP	HTTP/2	HTTP?
Pros	广泛使用、标准明确	Resource，数据格式灵活，易上手、架构风格	用途广泛、高性能	现代化、易上手、灵活、高性能	data structuring灵活、查询灵活和强大、现代化，schema一等公民
Cons	Heavy payload，复杂，难于上手，难于调试	过于理论化，标准不统一，服务接口没有强类型定义，低效的文本格式	流式数据处理能力欠缺，没有很好的发展起来	不支持HTTP/1.1，浏览器支持有待完善	上手困难，容易出现性能问题

Tubi Scala team & gRPC

- **Rainmaker** (广告)
- **Delphi** (推荐系统在线服务)
- **Data**
- **ML** (广告, 推荐)

```
commit 5a23800c926a7bd47d2f2fc5811078626e36449f
Author: Marios <marios@tubi.tv>
Date: Wed Feb 1 00:02:15 2017 -0800
```

initial commit, single proto definition

```
commit 5c607576603b7004f1e221075bfe4c83959dba39
Author: Marios <marios@tubi.tv>
Date: Thu Mar 2 01:59:40 2017 -0800
```

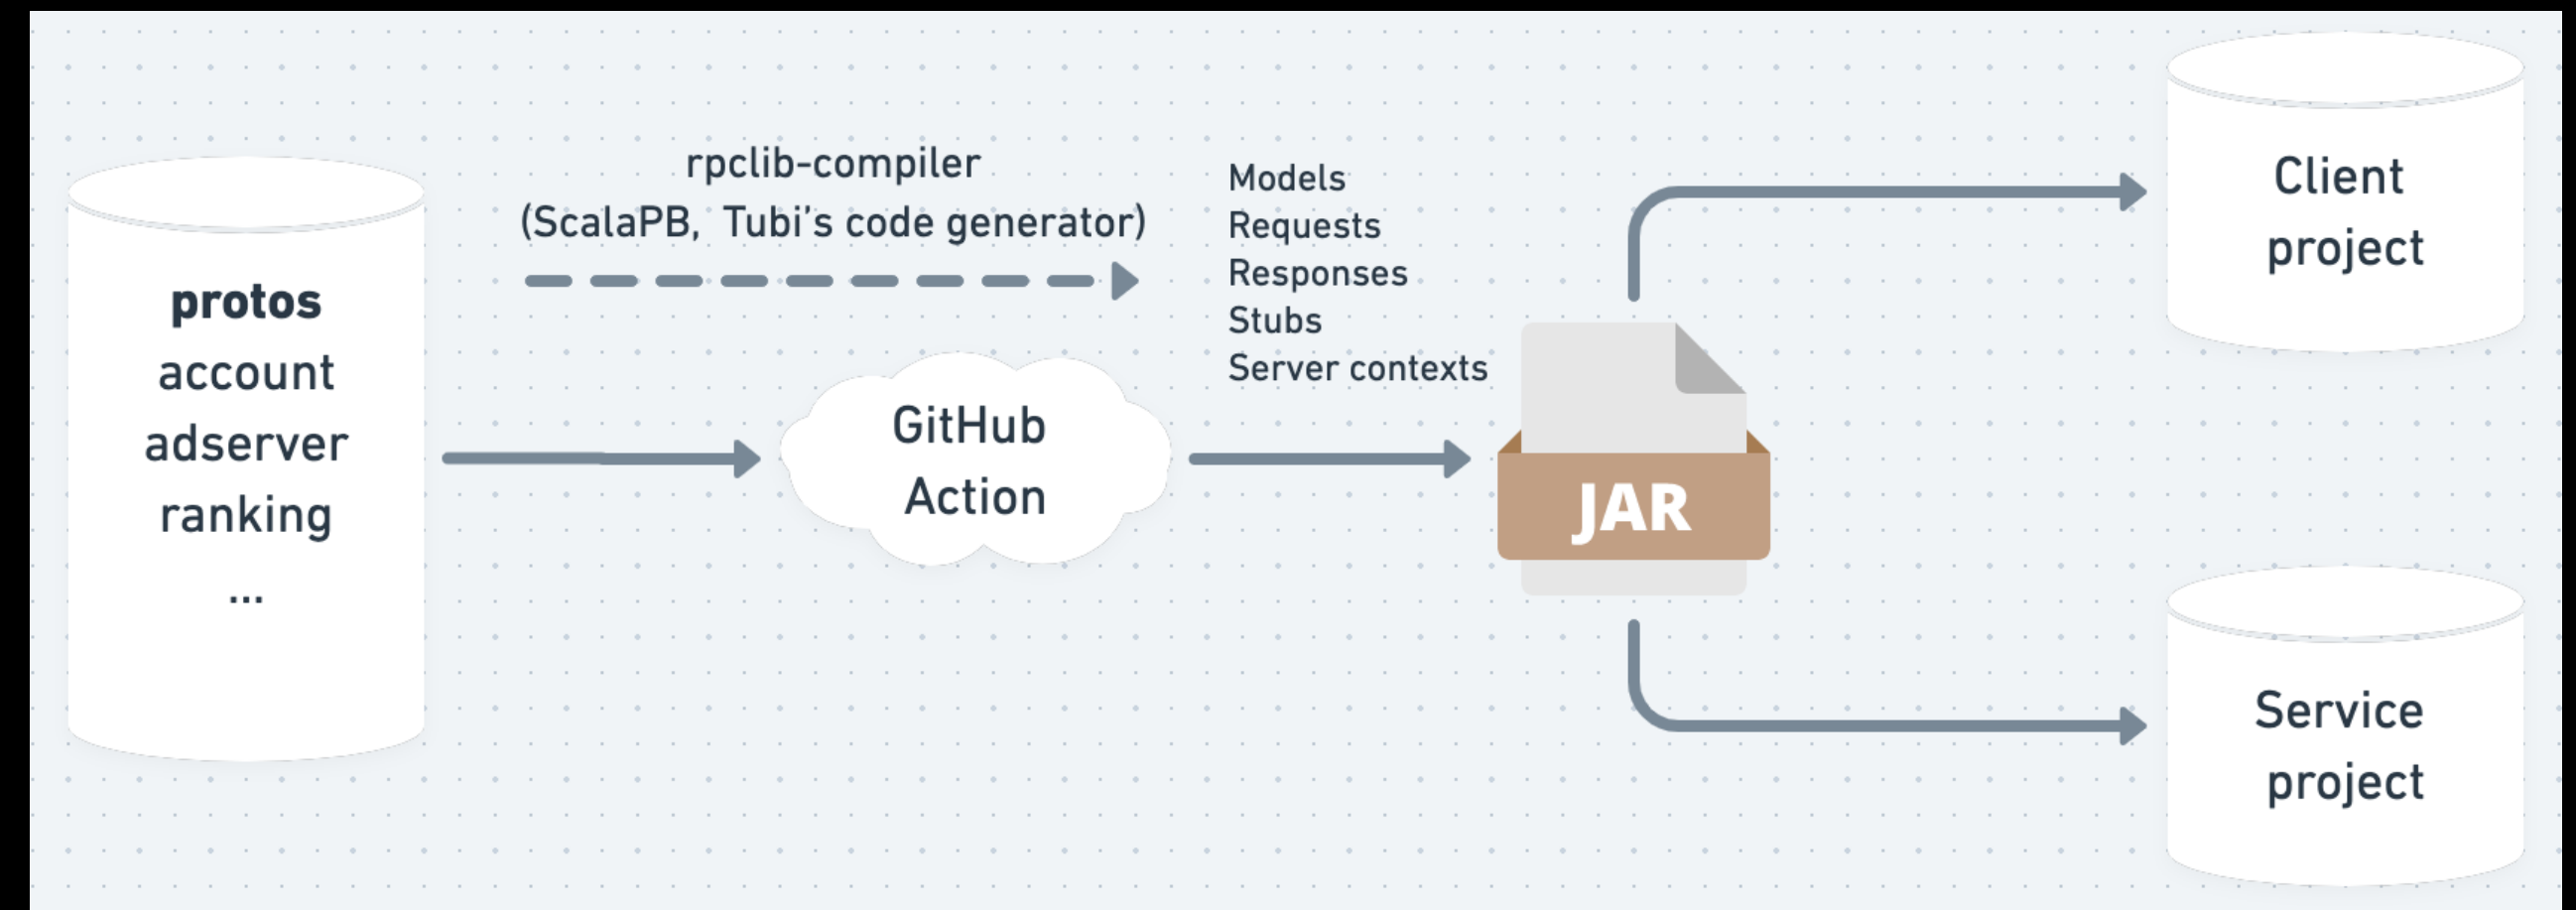
Some basic stubs to test with and bazel build

Why gRPC

- 面向服务定义而非资源定义，适合各种业务需求
- 强类型约束，降低开发和维护成本，提高工程师生产力
- 多语言支持，不同团队可以根据业务需要选择更合适的编程语言
- 不同团队可以通过集中的ProtoBuf代码库轻松的共享服务定义
- ProtoBuf高效的序列化性能和较小的序列化体积
- 原生支持 stream 操作, 大数据量传输更加高效
- 性能好，可支持高吞吐低延迟服务
- 社区活跃、生态快速成熟

开发流程

- protos
 - mono repo
 - 集中存储tubi内部各个gRPC服务定义
 - ScalaPB auxiliary options
- rpclib-compiler
 - tubi自研的sbt插件
 - 调用ScalaPB编译protobufs生成models类型定义(messages和enums)
 - 调用自研的code generator生成client stub和server skeleton
- rpclib-runtime
 - 服务器端/客户端辅助方法
 - MDC和Envoy配置(x-request-id/x-tubi-header)
 - 重试策略、超时配置、断路器策略



测试

- 单元测试和集成测试
- grpcurl和grpc_cli
- 本地和staging环境验证
- grpc-ui

gRPC Server Target

ranking:50051

Use TLS

Restart Connection

Use local proto

Use request metadata

	Key	Value
<div><div></div></div>	x-request-id	test

+

Services

personal.PersonalizeMatrix

Methods

PersonalizeAutoplay

```
{
  "request_context": {
    "device_id": "abc",
    "user_id": 0,
    "platform": 1,
    "is_authenticated": false,
    "ip_address": "127.0.0.1",
    "utc_timestamp": "2017-01-15T01:30:15.01Z",
    "country": 1,
    "timezone": "San Francisco"
  },
  "current_content": {
    "id": 123201
  }
}
```

Schema Input

```
message PersonalizeAutoplayRequest {
  .personal.RequestContext request_context = 1;
  .personal.Content source_content = 2 [(.personal.Content) = 3];
  oneof candidates {
    .personal.Episodes episodes = 4 [(.personal.Episodes) = 5];
    .personal.Container container = 5 [(.personal.Container) = 6];
    .personal.Sequels sequels = 10 [(.personal.Sequels) = 11];
  }
}
```

SUBMIT

Response:

```
{
  "content": [
    {
      "id": 9999999999,
      "contentType": "SERIES"
    },
  ],
}
```

Time : 112ns

部署

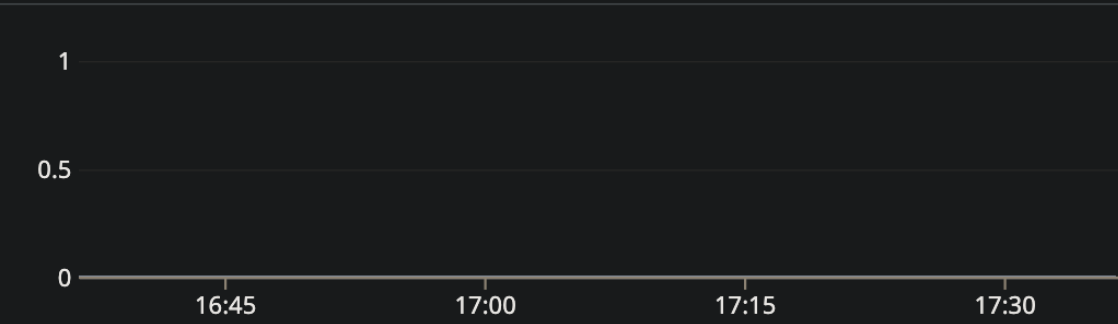
- K8s紧密集成
 - gRPC Health Check
 - 服务发现 -> k8s service
 - k8s readiness和liveness
- Envoy

```
readinessProbe:
  initialDelaySeconds: 20
  timeoutSeconds: 5
  exec:
    command: [ "/bin/grpc_health_probe", "-addr=:50051", "-service=ranking" ]
livenessProbe:
  initialDelaySeconds: 20
  timeoutSeconds: 5
  exec:
    command: [ "/bin/grpc_health_probe", "-addr=:50051", "-service=ranking" ]
```

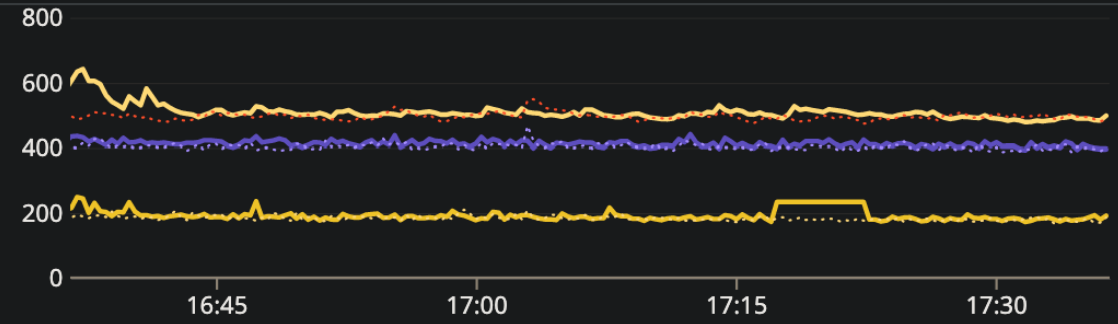
Metrics

- 系统 Metrics
- App 通用 Metrics
- App metrics

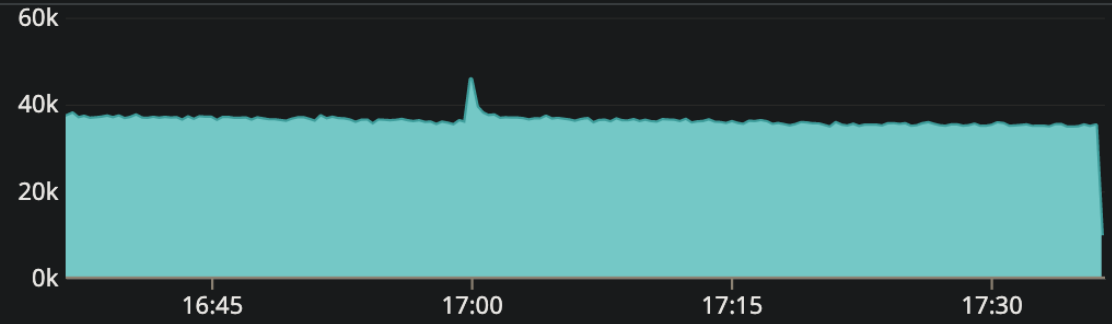
Network Errors



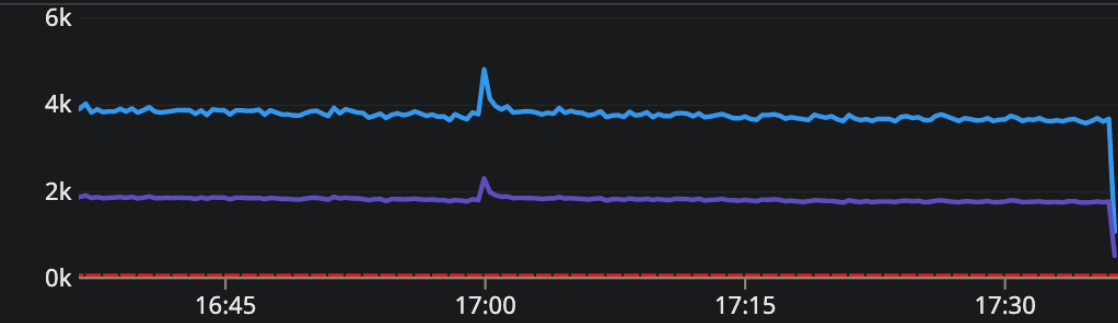
Envoy: Service Performance



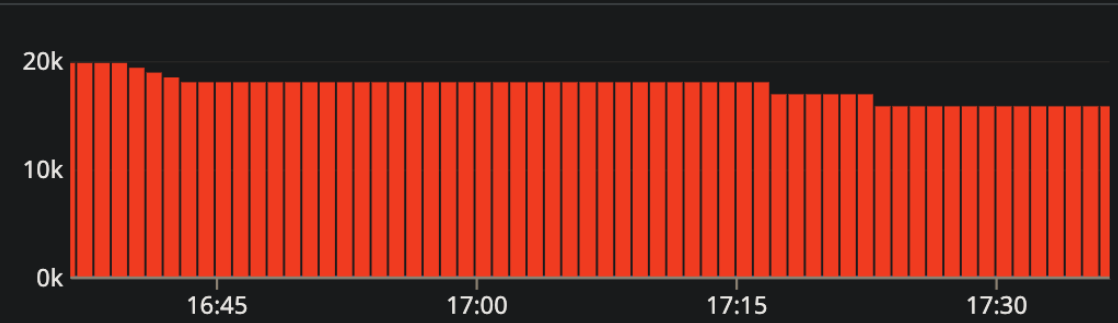
Envoy: Response Class Distribution



Envoy: Service Total & Routed RPS (lines) vs Total Active Count (b...



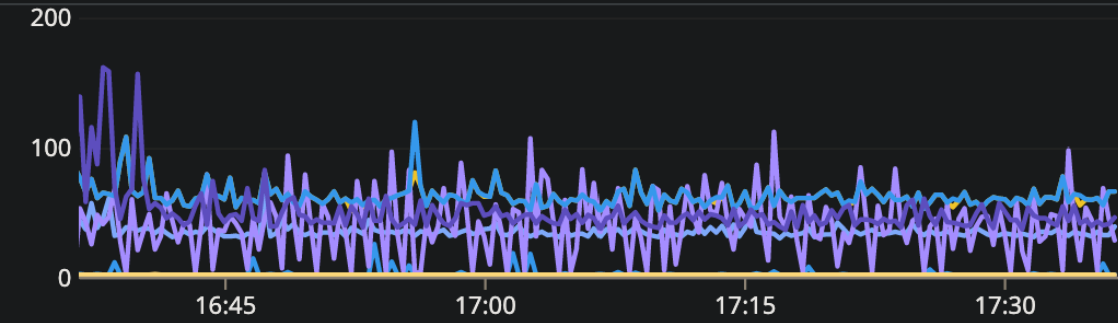
Envoy: Service Total & Routed CPS (lines) vs Total Active Count (b...



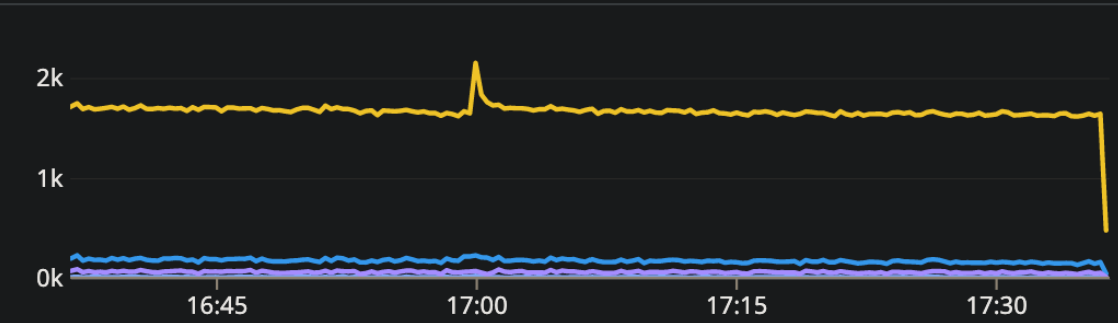
Envoy: Service Non-2xx Responses Count Breakdown (vs total o...



Envoy: Dependencies Request Time (in ms)



Envoy: Dependencies RPS



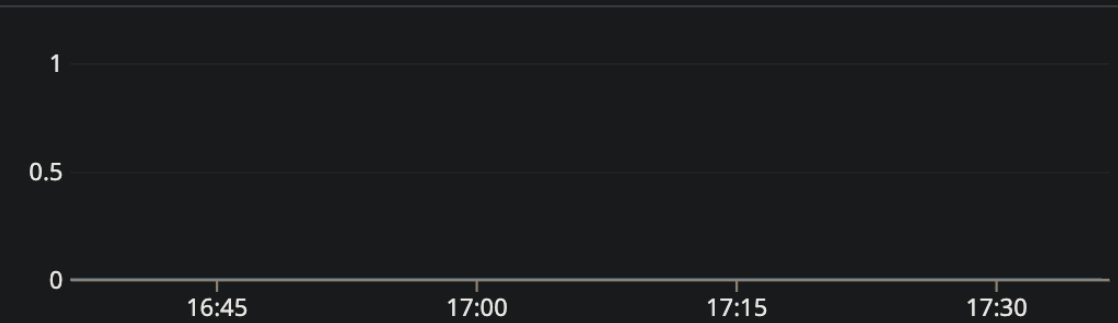
Envoy: Dependencies Non-2xx Responses Count Breakdown



Envoy: Dependencies Retry Total (positive bars) vs Failed (negati...



Envoy: Upstream Requests

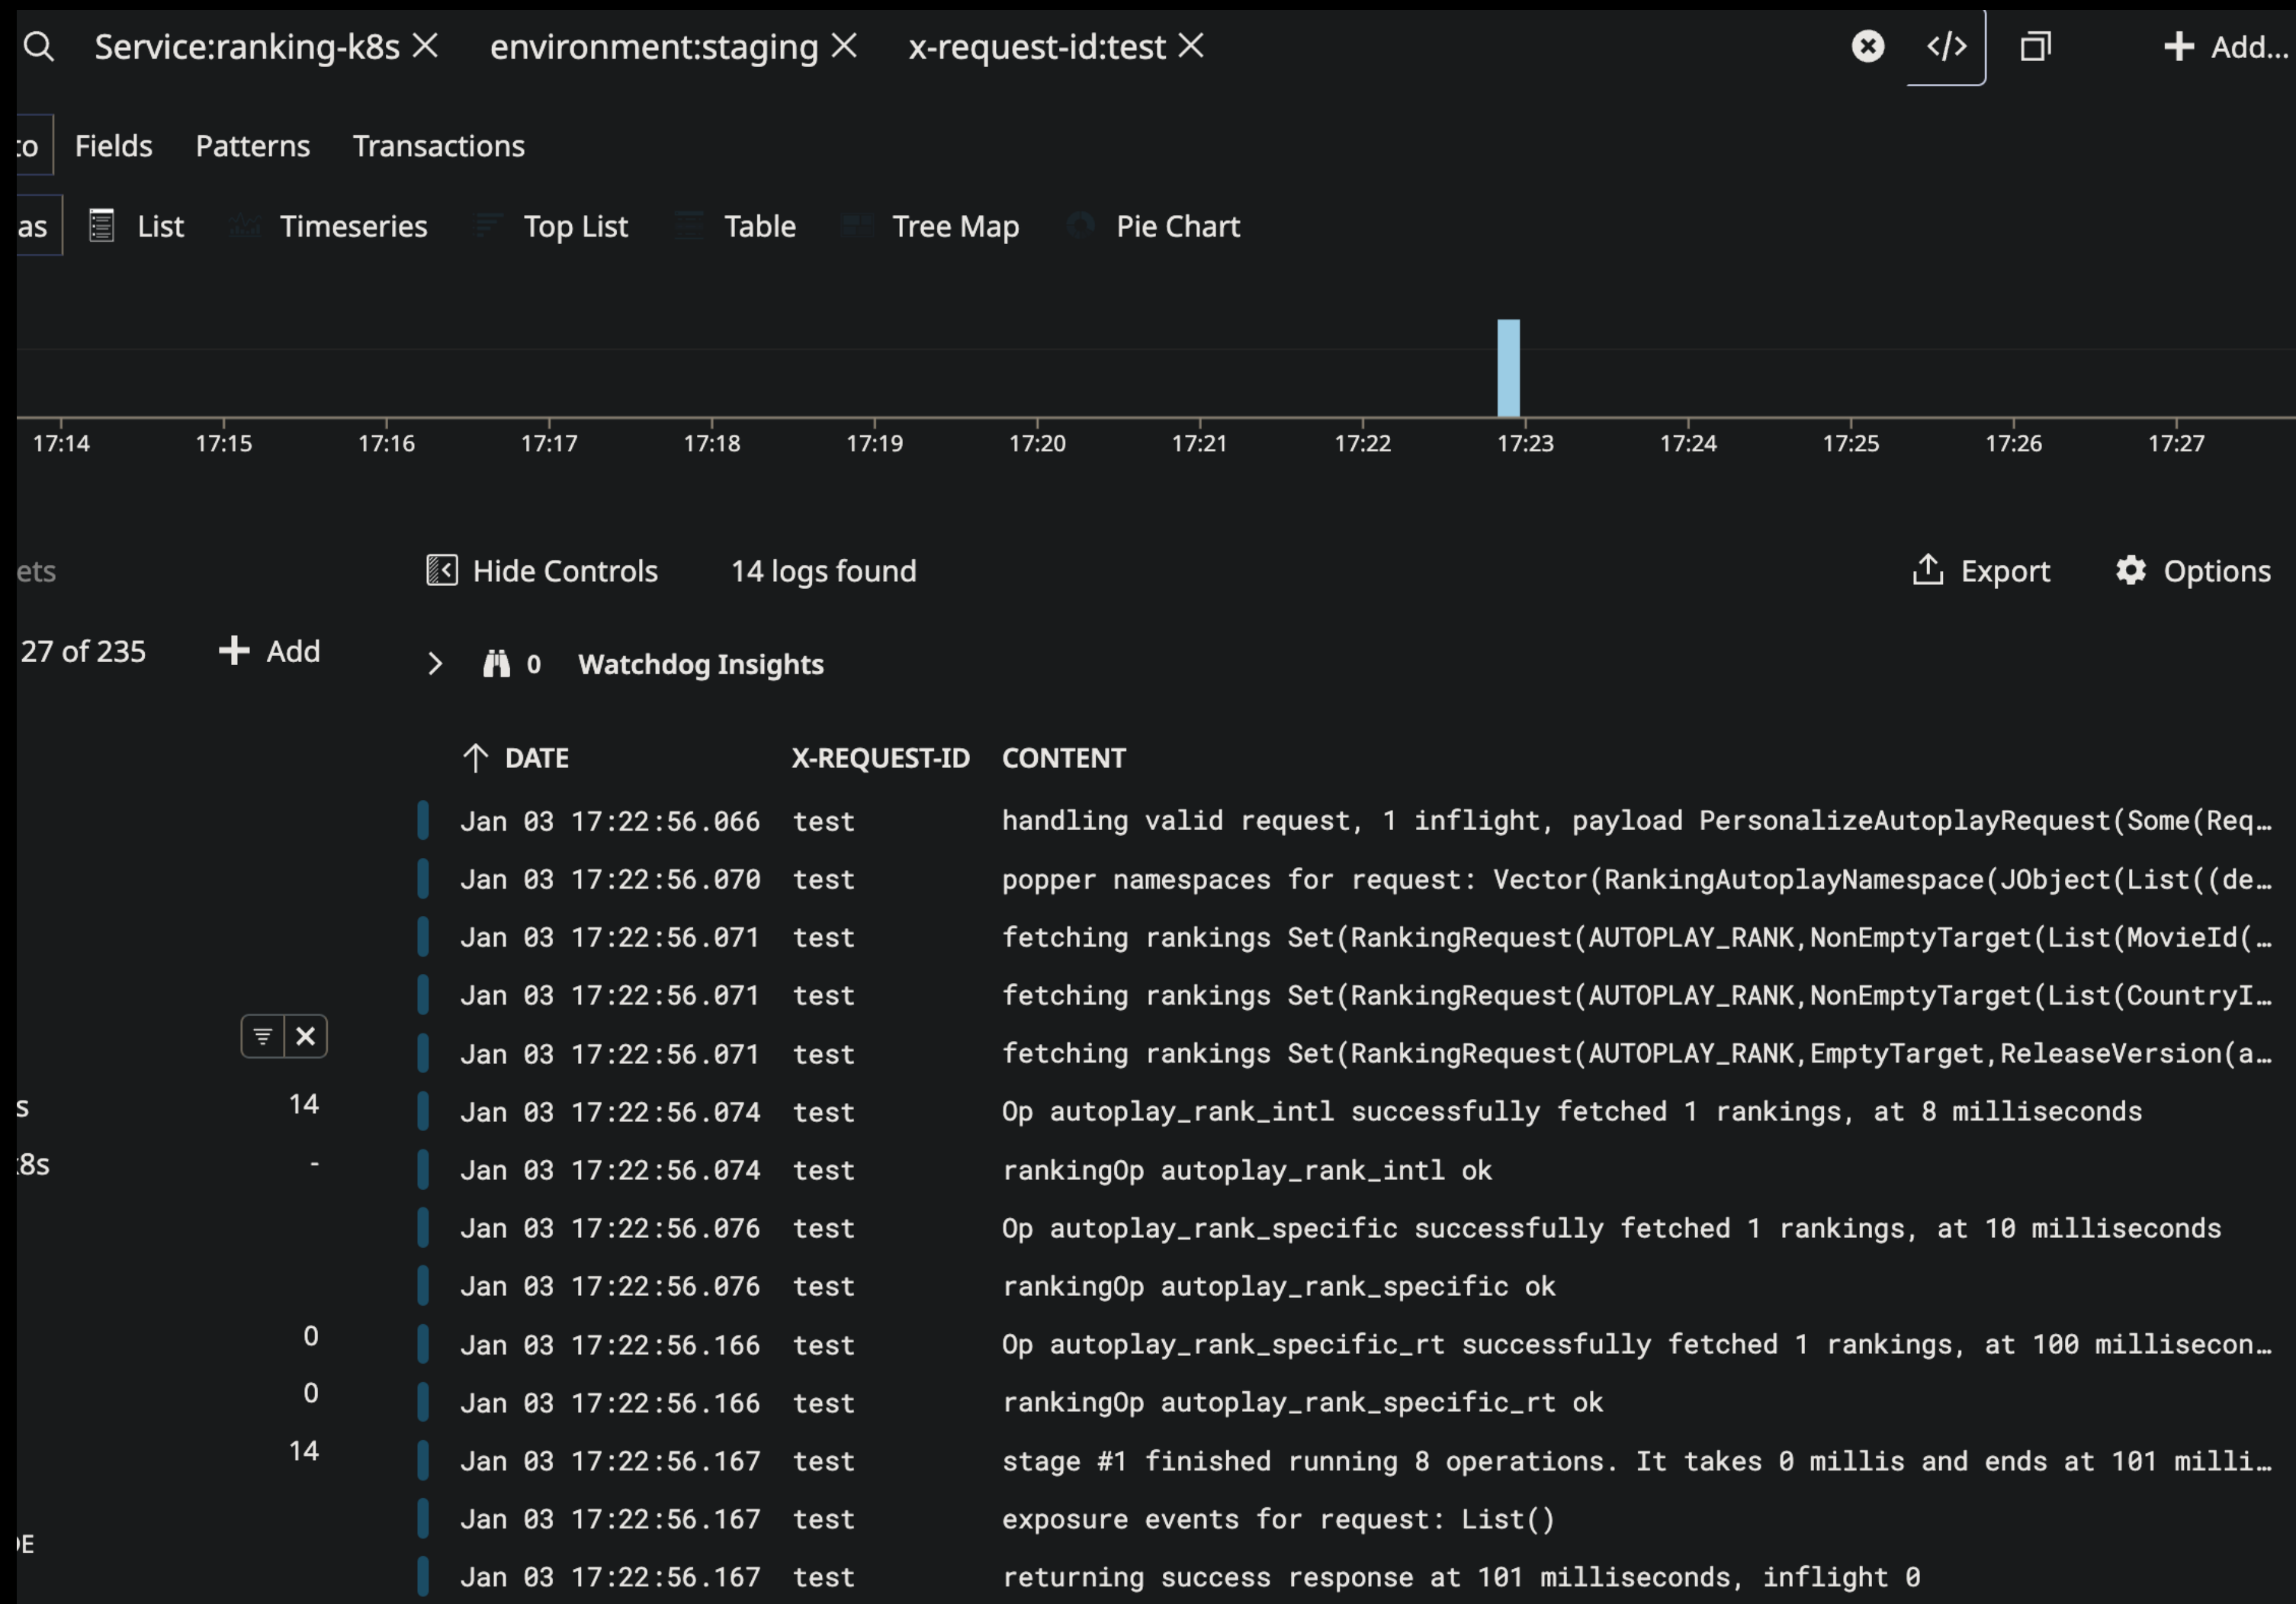


Envoy: Return Non-2xx Responses to Upstream Count Breakdown



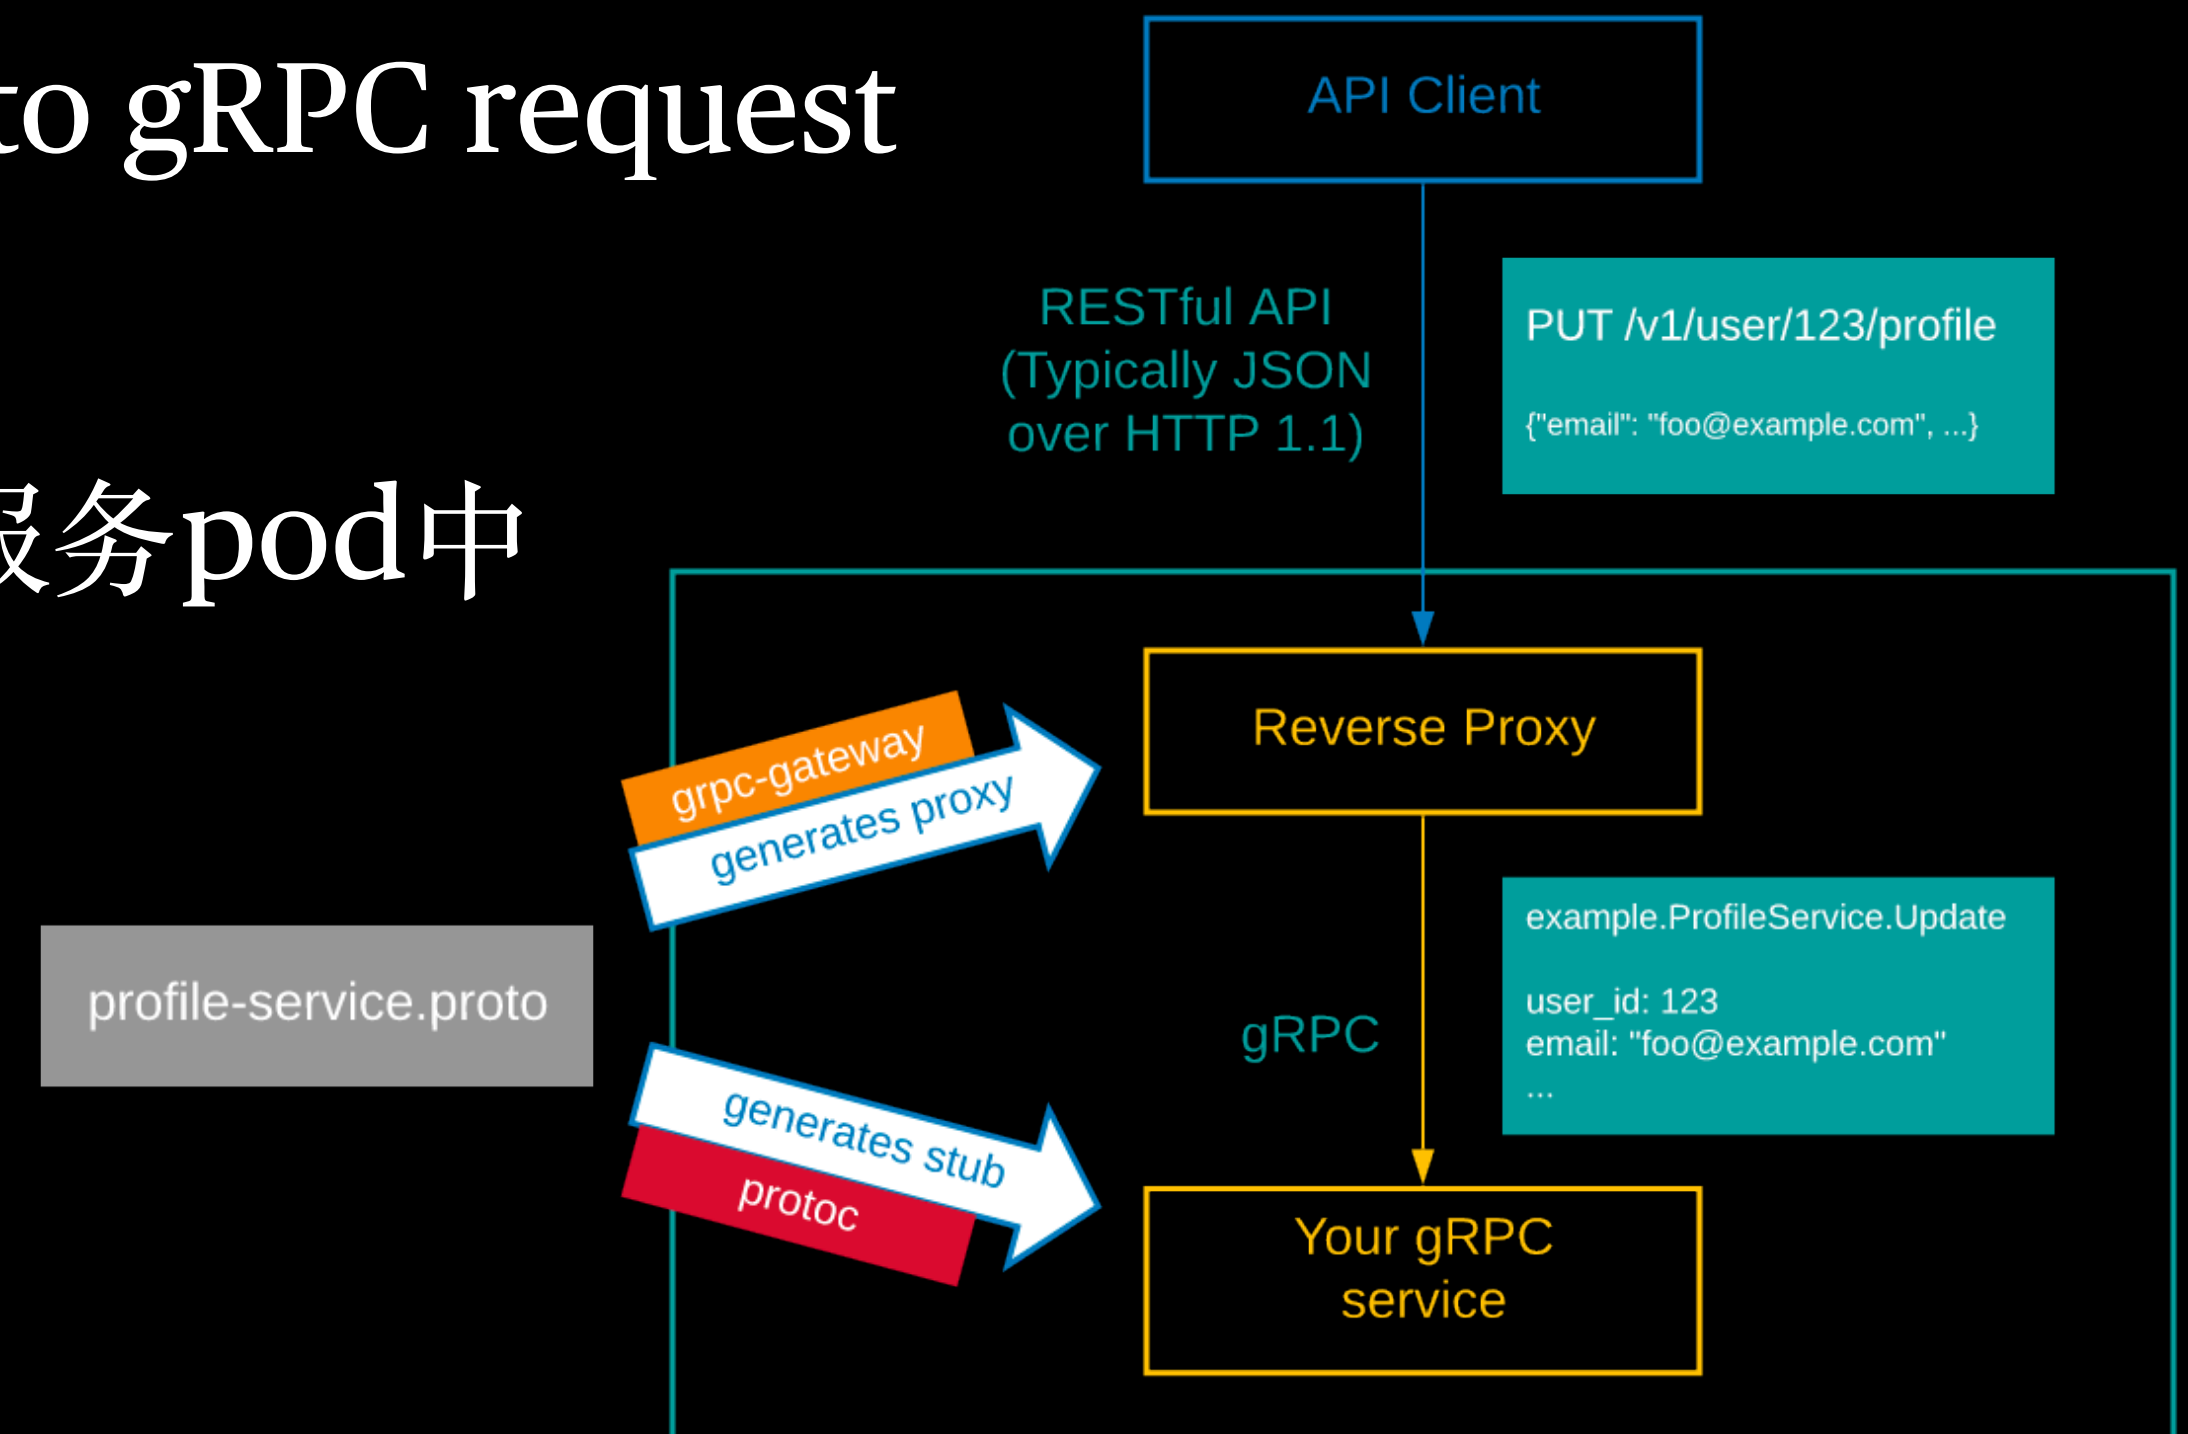
Logging with MDC

- Mapped Diagnostic Context
- x-request-id
- 客户端id、客户端ip、客户端参数
- 在海量日志数据中快速捞到可用的日志信息



REST on gRPC

- grpc-bridge
 - Internal service proxy HTTP request to gRPC request
 - 基于grpc-gateway
 - 以sidecar container方式注入到gRPC服务pod中



ref: [architecture_introduction_diagram](#)

Akka-gRPC迁移

感谢Weiwen同学的迁移工作

Akka-gRPC 概述

- 基于Akka Streams和Akka HTTP (http2)
- 流式gRPC服务
- 代码生成器
 - Model类定义
 - 基于Akka Stream Sources 的服务端API skeleton
 - Akka HTTP route用来辅助在Akka http上serving gRPC
 - Client stub
- gRPC runtime实现
- gRPC-web支持

Why Akka-gRPC

- Tubi后端重度依赖Akka生态
- 降低维护现有rpclib-runtime的成本
 - 底层实现绑定低版本ScalaPB，间接导致无法升级grpc-api/grpc-core
- Ready for production
- 高性能
 - 官方benchmark
- 更好的特性支持，如server reflection、gRPC-web、metrics

Akka-gRPC性能

- Intel(R) Xeon(R) Platinum 8259CL CPU @ 2.50GHz
- OS: Ubuntu 20.04.4 LTS
- Naïve random UUID generation

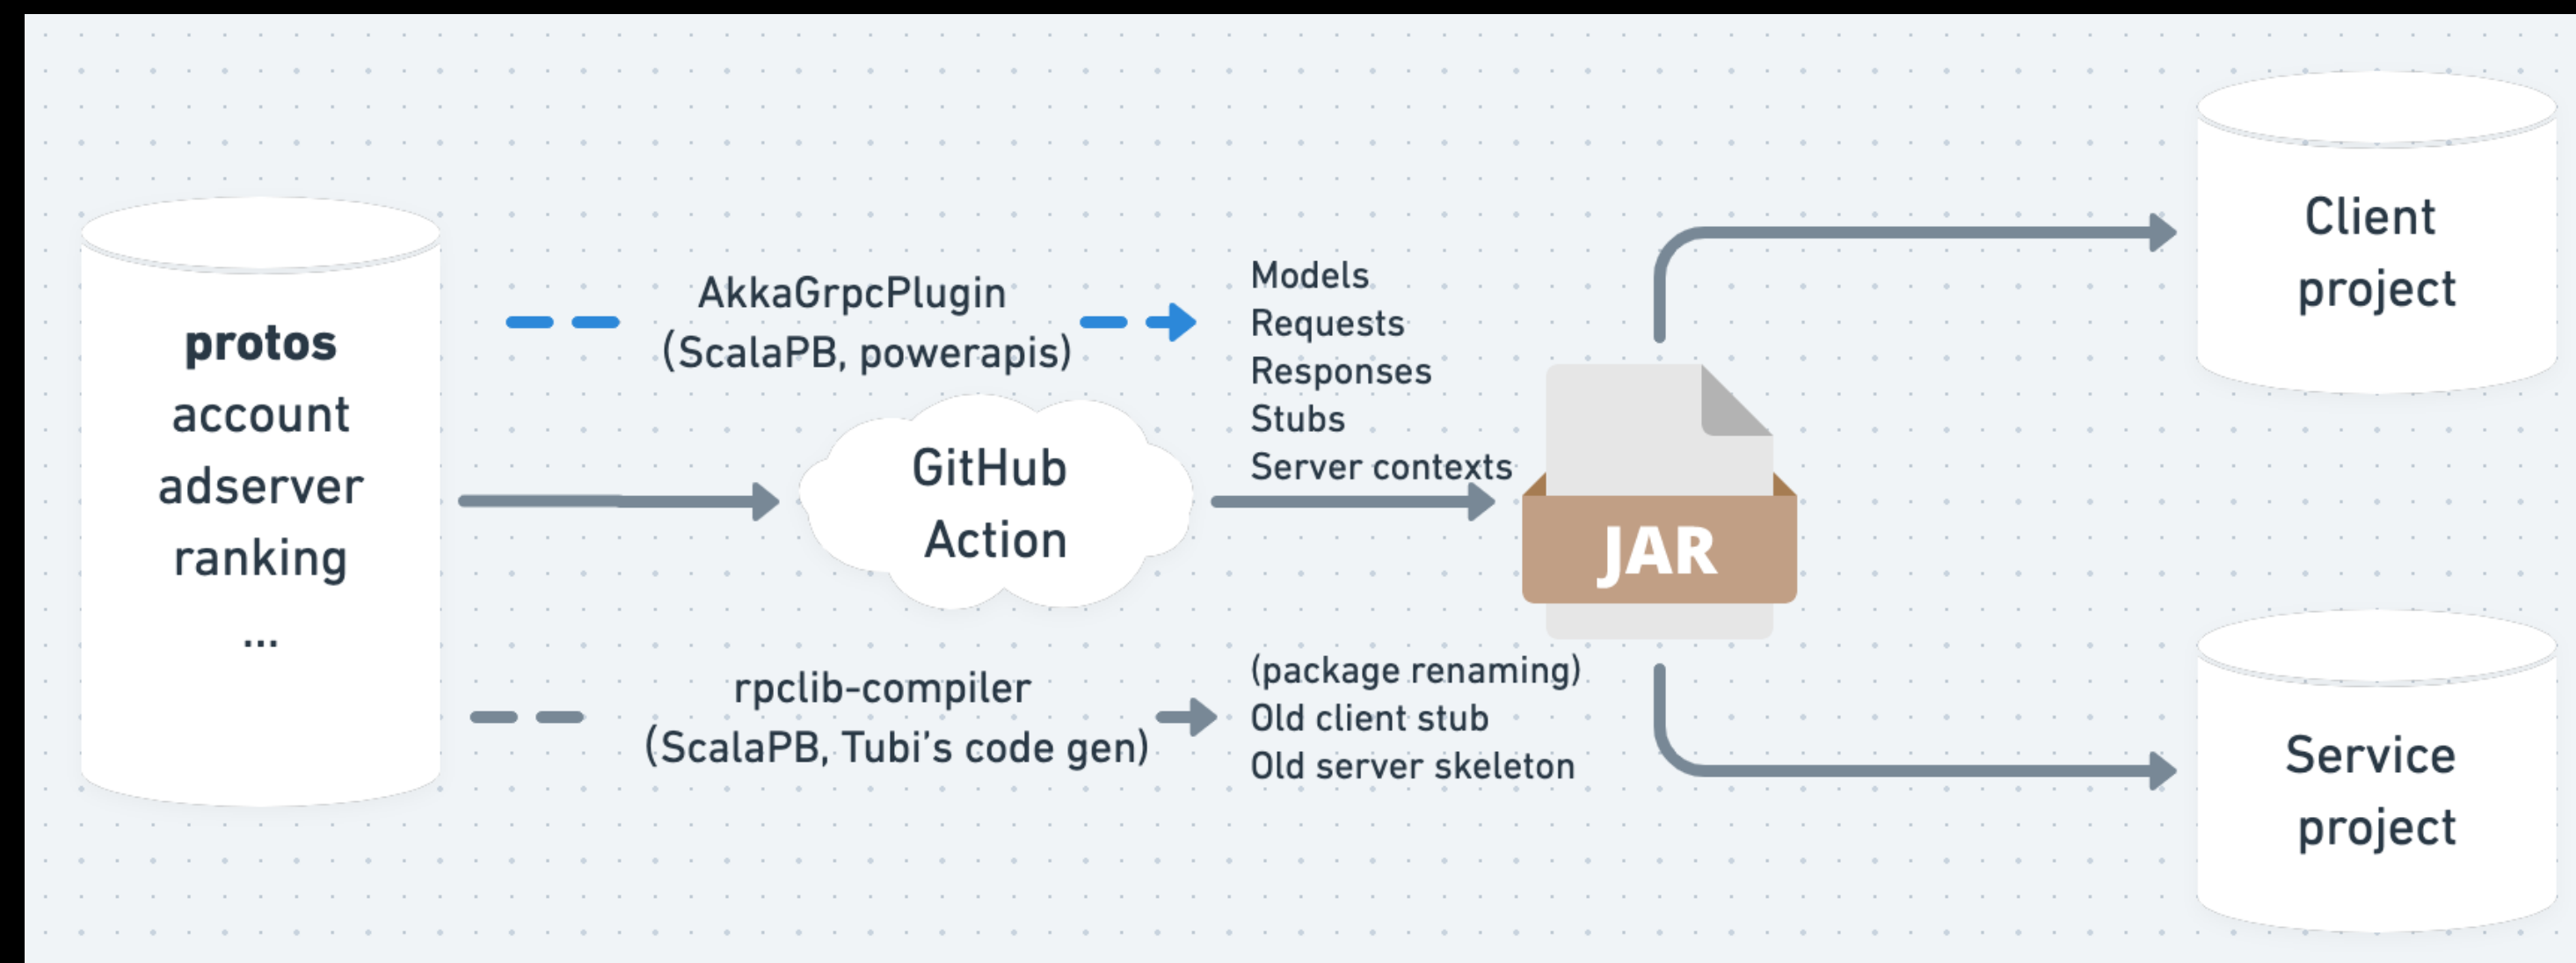
		Akka-grpc(1.1.0)	Rpclib(2.0.16)	Akka-grpc(2.1.4)	Rpclib(3.0.0)
Unary	Summary	<div>1Count: 200000</div> <div>2Total: 51.01 s</div> <div>3Slowest: 25.18 ms</div> <div>4Fastest: 0.42 ms</div> <div>5Average: 4.68 ms</div> <div>6Requests/sec: 3921.14</div>	<div>1Count: 200000</div> <div>2Total: 58.23 s</div> <div>3Slowest: 824.90 ms</div> <div>4Fastest: 0.29 ms</div> <div>5Average: 5.68 ms</div> <div>6Requests/sec: 3434.54</div>	<div>1Count: 200000</div> <div>2Total: 39.00 s</div> <div>3Slowest: 26.68 ms</div> <div>4Fastest: 0.29 ms</div> <div>5Average: 3.37 ms</div> <div>6Requests/sec: 5128.60</div>	<div>1Count: 200000</div> <div>2Total: 42.21 s</div> <div>3Slowest: 39.53 ms</div> <div>4Fastest: 0.26 ms</div> <div>5Average: 3.61 ms</div> <div>6Requests/sec: 4737.99</div>
	Latency distribution	<div>110 % in 2.05 ms</div> <div>225 % in 3.21 ms</div> <div>350 % in 4.42 ms</div> <div>475 % in 5.76 ms</div> <div>590 % in 7.69 ms</div> <div>695 % in 8.90 ms</div> <div>799 % in 11.10 ms</div>	<div>110 % in 1.88 ms</div> <div>225 % in 2.89 ms</div> <div>350 % in 4.14 ms</div> <div>475 % in 6.50 ms</div> <div>590 % in 10.27 ms</div> <div>695 % in 13.85 ms</div> <div>799 % in 26.75 ms</div>	<div>110 % in 1.40 ms</div> <div>225 % in 2.12 ms</div> <div>350 % in 2.97 ms</div> <div>475 % in 4.17 ms</div> <div>590 % in 6.06 ms</div> <div>695 % in 7.15 ms</div> <div>799 % in 8.99 ms</div>	<div>110 % in 1.41 ms</div> <div>225 % in 2.24 ms</div> <div>350 % in 3.23 ms</div> <div>475 % in 4.46 ms</div> <div>590 % in 6.39 ms</div> <div>695 % in 7.64 ms</div> <div>799 % in 10.04 ms</div>
Client Stream	Summary	<div>1Count: 100000</div> <div>2Total: 62.33 s</div> <div>3Slowest: 58.52 ms</div> <div>4Fastest: 0.82 ms</div> <div>5Average: 13.72 ms</div> <div>6Requests/sec: 1604.46</div>	<div>1Count: 100000</div> <div>2Total: 200.84 s</div> <div>3Slowest: 182.31 ms</div> <div>4Fastest: 3.02 ms</div> <div>5Average: 59.12 ms</div> <div>6Requests/sec: 497.90</div>	<div>1Count: 100000</div> <div>2Total: 53.92 s</div> <div>3Slowest: 63.82 ms</div> <div>4Fastest: 0.65 ms</div> <div>5Average: 11.55 ms</div> <div>6Requests/sec: 1854.61</div>	<div>1Count: 100000</div> <div>2Total: 185.94 s</div> <div>3Slowest: 113.62 ms</div> <div>4Fastest: 4.41 ms</div> <div>5Average: 54.42 ms</div> <div>6Requests/sec: 537.81</div>
	Latency distribution	<div>110 % in 7.53 ms</div> <div>225 % in 10.63 ms</div> <div>350 % in 13.52 ms</div> <div>475 % in 16.75 ms</div> <div>590 % in 20.13 ms</div> <div>695 % in 22.09 ms</div> <div>799 % in 26.41 ms</div>	<div>110 % in 48.74 ms</div> <div>225 % in 54.06 ms</div> <div>350 % in 58.68 ms</div> <div>475 % in 63.71 ms</div> <div>590 % in 69.67 ms</div> <div>695 % in 74.76 ms</div> <div>799 % in 90.67 ms</div>	<div>110 % in 6.56 ms</div> <div>225 % in 8.65 ms</div> <div>350 % in 10.98 ms</div> <div>475 % in 14.08 ms</div> <div>590 % in 17.47 ms</div> <div>695 % in 19.39 ms</div> <div>799 % in 23.78 ms</div>	<div>110 % in 44.57 ms</div> <div>225 % in 50.02 ms</div> <div>350 % in 54.66 ms</div> <div>475 % in 59.49 ms</div> <div>590 % in 64.40 ms</div> <div>695 % in 67.81 ms</div> <div>799 % in 75.93 ms</div>

Akka-gRPC性能(cont.)

		Akka-grpc(1.1.0)	Rpclib(2.0.16)	Akka-grpc(2.1.4)	Rpclib(3.0.0)
Server Stream	Summary	<div><div>1</div><div>Count: 100000</div></div> <div><div>2</div><div>Total: 64.84 s</div></div> <div><div>3</div><div>Slowest: 131.21 ms</div></div> <div><div>4</div><div>Fastest: 1.06 ms</div></div> <div><div>5</div><div>Average: 14.84 ms</div></div> <div><div>6</div><div>Requests/sec: 1542.33</div></div>	<div><div>1</div><div>Count: 100000</div></div> <div><div>2</div><div>Total: 97.66 s</div></div> <div><div>3</div><div>Slowest: 121.97 ms</div></div> <div><div>4</div><div>Fastest: 1.45 ms</div></div> <div><div>5</div><div>Average: 24.62 ms</div></div> <div><div>6</div><div>Requests/sec: 1024.00</div></div>	<div><div>1</div><div>Count: 100000</div></div> <div><div>2</div><div>Total: 64.78 s</div></div> <div><div>3</div><div>Slowest: 75.62 ms</div></div> <div><div>4</div><div>Fastest: 1.16 ms</div></div> <div><div>5</div><div>Average: 14.99 ms</div></div> <div><div>6</div><div>Requests/sec: 1543.76</div></div>	<div><div>1</div><div>Count: 100000</div></div> <div><div>2</div><div>Total: 81.64 s</div></div> <div><div>3</div><div>Slowest: 82.76 ms</div></div> <div><div>4</div><div>Fastest: 1.11 ms</div></div> <div><div>5</div><div>Average: 20.48 ms</div></div> <div><div>6</div><div>Requests/sec: 1224.83</div></div>
	Latency distribution	<div><div>1</div><div>10 % in 8.69 ms</div></div> <div><div>2</div><div>25 % in 11.61 ms</div></div> <div><div>3</div><div>50 % in 14.69 ms</div></div> <div><div>4</div><div>75 % in 17.82 ms</div></div> <div><div>5</div><div>90 % in 20.88 ms</div></div> <div><div>6</div><div>95 % in 22.89 ms</div></div> <div><div>7</div><div>99 % in 27.74 ms</div></div>	<div><div>1</div><div>10 % in 13.67 ms</div></div> <div><div>2</div><div>25 % in 18.12 ms</div></div> <div><div>3</div><div>50 % in 23.50 ms</div></div> <div><div>4</div><div>75 % in 29.59 ms</div></div> <div><div>5</div><div>90 % in 36.38 ms</div></div> <div><div>6</div><div>95 % in 41.78 ms</div></div> <div><div>7</div><div>99 % in 55.01 ms</div></div>	<div><div>1</div><div>10 % in 9.26 ms</div></div> <div><div>2</div><div>25 % in 11.88 ms</div></div> <div><div>3</div><div>50 % in 14.69 ms</div></div> <div><div>4</div><div>75 % in 17.88 ms</div></div> <div><div>5</div><div>90 % in 21.04 ms</div></div> <div><div>6</div><div>95 % in 23.13 ms</div></div> <div><div>7</div><div>99 % in 28.01 ms</div></div>	<div><div>1</div><div>10 % in 11.74 ms</div></div> <div><div>2</div><div>25 % in 15.42 ms</div></div> <div><div>3</div><div>50 % in 19.90 ms</div></div> <div><div>4</div><div>75 % in 24.92 ms</div></div> <div><div>5</div><div>90 % in 29.75 ms</div></div> <div><div>6</div><div>95 % in 32.97 ms</div></div> <div><div>7</div><div>99 % in 39.91 ms</div></div>
Bidi Stream	Summary	<div><div>1</div><div>Count: 100000</div></div> <div><div>2</div><div>Total: 109.99 s</div></div> <div><div>3</div><div>Slowest: 107.71 ms</div></div> <div><div>4</div><div>Fastest: 1.73 ms</div></div> <div><div>5</div><div>Average: 27.06 ms</div></div> <div><div>6</div><div>Requests/sec: 909.18</div></div>	<div><div>1</div><div>Count: 100000</div></div> <div><div>2</div><div>Total: 309.82 s</div></div> <div><div>3</div><div>Slowest: 618.19 ms</div></div> <div><div>4</div><div>Fastest: 6.45 ms</div></div> <div><div>5</div><div>Average: 87.64 ms</div></div> <div><div>6</div><div>Requests/sec: 322.76</div></div>	<div><div>1</div><div>Count: 100000</div></div> <div><div>2</div><div>Total: 97.59 s</div></div> <div><div>3</div><div>Slowest: 108.56 ms</div></div> <div><div>4</div><div>Fastest: 1.94 ms</div></div> <div><div>5</div><div>Average: 23.89 ms</div></div> <div><div>6</div><div>Requests/sec: 1024.71</div></div>	<div><div>1</div><div>Count: 100000</div></div> <div><div>2</div><div>Total: 275.75 s</div></div> <div><div>3</div><div>Slowest: 254.31 ms</div></div> <div><div>4</div><div>Fastest: 5.00 ms</div></div> <div><div>5</div><div>Average: 77.73 ms</div></div> <div><div>6</div><div>Requests/sec: 362.65</div></div>
	Latency distribution	<div><div>1</div><div>10 % in 15.04 ms</div></div> <div><div>2</div><div>25 % in 21.36 ms</div></div> <div><div>3</div><div>50 % in 27.53 ms</div></div> <div><div>4</div><div>75 % in 32.85 ms</div></div> <div><div>5</div><div>90 % in 37.71 ms</div></div> <div><div>6</div><div>95 % in 40.89 ms</div></div> <div><div>7</div><div>99 % in 48.30 ms</div></div>	<div><div>1</div><div>10 % in 66.18 ms</div></div> <div><div>2</div><div>25 % in 81.27 ms</div></div> <div><div>3</div><div>50 % in 90.16 ms</div></div> <div><div>4</div><div>75 % in 97.72 ms</div></div> <div><div>5</div><div>90 % in 104.81 ms</div></div> <div><div>6</div><div>95 % in 109.87 ms</div></div> <div><div>7</div><div>99 % in 121.93 ms</div></div>	<div><div>1</div><div>10 % in 14.09 ms</div></div> <div><div>2</div><div>25 % in 18.94 ms</div></div> <div><div>3</div><div>50 % in 23.79 ms</div></div> <div><div>4</div><div>75 % in 28.73 ms</div></div> <div><div>5</div><div>90 % in 33.55 ms</div></div> <div><div>6</div><div>95 % in 36.65 ms</div></div> <div><div>7</div><div>99 % in 43.09 ms</div></div>	<div><div>1</div><div>10 % in 57.39 ms</div></div> <div><div>2</div><div>25 % in 71.09 ms</div></div> <div><div>3</div><div>50 % in 79.63 ms</div></div> <div><div>4</div><div>75 % in 86.98 ms</div></div> <div><div>5</div><div>90 % in 94.15 ms</div></div> <div><div>6</div><div>95 % in 99.78 ms</div></div> <div><div>7</div><div>99 % in 117.92 ms</div></div>

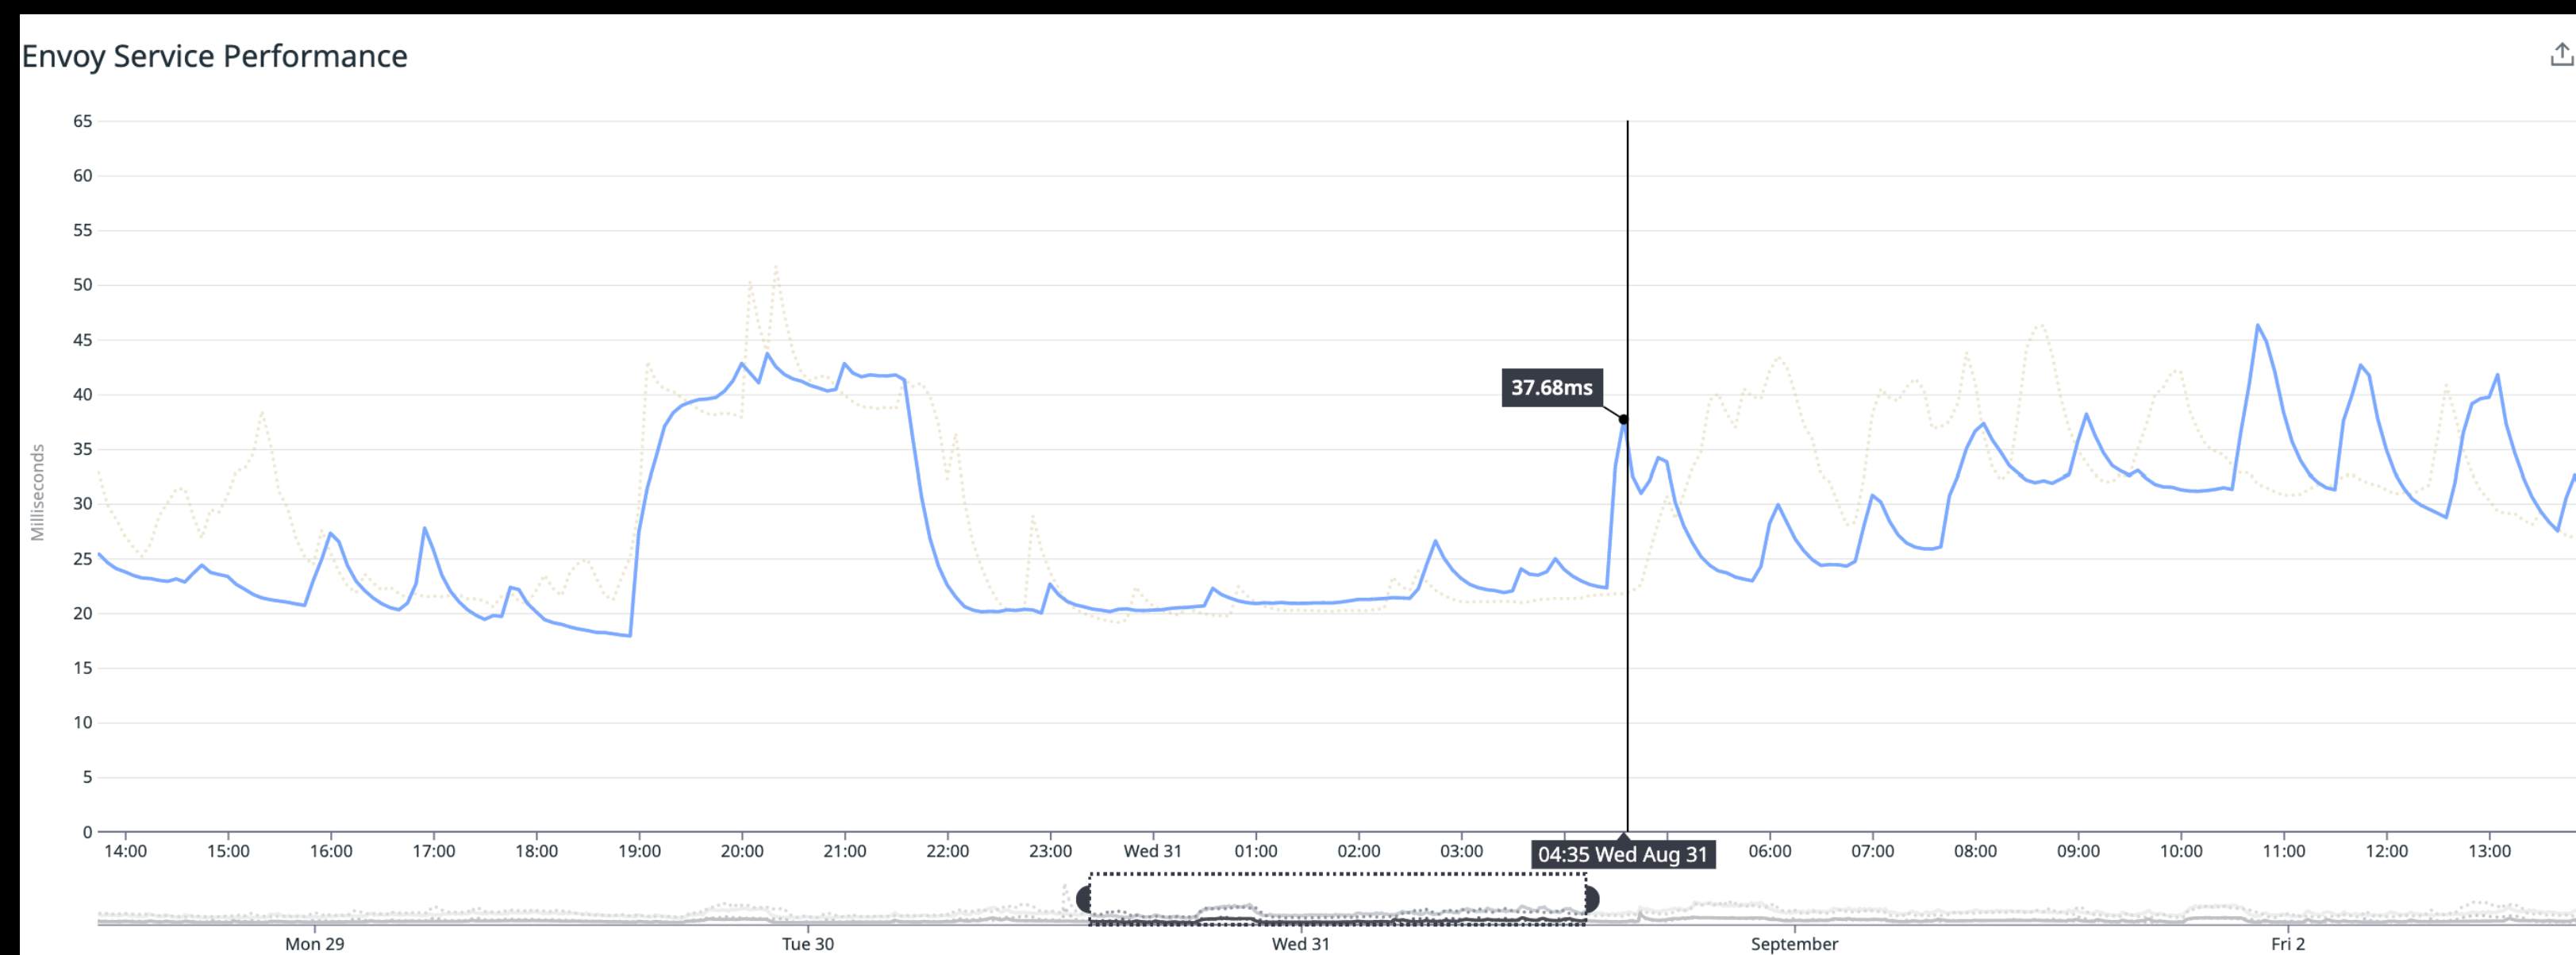
迁移步骤

- 更新rpclib-runtime使用Akka-gRPC
- 更新protos使用Akka-gRPC sbt plugin
- 利用Akka-gRPC的power apis重新实现MDC、interceptor机制
- 难点在于兼容现有老的代码
 - ScalaPB / rpclib-compiler / Akka-gRPC
 - 通过class renaming共存
 - 提供一个大版本的过渡期



结果

- 性能指标维持不变
 - latency基本维持不变
 - 内存和cpu基本维持不变
- 可维护性提升
 - 升级grpc依赖到新版
 - 代码可读性提高
 - Akka社区力量



Q&A and Thanks!